

Week 2 (Jan 11 - 15)

Process

- Process - Program in execution
 - Has address space the program can use
 - Has State (registers, including program counter and stack pointer)
- OS tracks processes in process table
- Processes can create other processes
 - Process tree tracks these relationships

Process = Address Space + Thread of Control

Inside a UNIX Process

- Three Segments:
 - Text: contains the program code
 - Data: contains the program data
 - * static variables
 - * malloc, new
 - stack:
 - * Automatic variables
 - * Procedure call info
- Address space Growth
 - Text: No growth
 - Data and Stack grow towards each other

Deadlocks are caused by Circular dependencies

Interprocess Communication

- Processes work to exchange info with each other
- Many ways
 - Network
 - Pipes

System Calls

- OS runs in priveleged (supervisor) mode
 - User Processes do not
- Programs want OS to perform service
 - Access a file
 - Create process
- Accomplished with System Calls

How System Calls Work

- User program enters priveleged mode through well defined entrypoint
- Program passes relevant info into OS
- OS performs service
 - Able to do so
 - Service is permitted for user process
- Example system calls
 - fork
 - waitpid
 - execvp

Operating System Types

- Monolithic
 - All just a program
 - * Any piece has access to other pieces
 - Sometimes Modular
 - * Extra pieces dynamically added
 - * Extra pieces part of whole
- Layered
 - CPU Supports layers of priveledge
 - * Intel supports 4
 - * User at outermost layer
 - Use priveledge layers to enforce separation and provide security
 - * Outer layers can only move inward at certain gates
 - Not widely used, but a good idea
- Micro-Kernels
 - Bare minimum in kernel
 - Everything else in User-space
 - Inefficient

Scheduling

- How are processes scheduled?
 - When they enter system
 - * Batch Systems
 - New job runs scheduler
 - Scheduling done when job voluntarily gives up time
 - Fixed interval scheduling
 - * Necessary for interactive systems
- Goals of Scheduling
 - Fairness: Fair chance at CPU
 - Enforcement: Ensure policy is carried out
 - Balance: Keep all parts of system busy

- Batch
 - * Throughput
 - Work/time
 - * Turnaround time
 - Like response time, but for batch jobs
 - * CPU utilization
- Interactive
 - * Response time
 - Can measure avg or other metrics
 - Good to measure # times faster than certain time (measure # of outliers)
 - * Proportionality: meet user expectations
- Real time
 - * meet deadlines
 - * predictability

Types of Schedulers

- Batch Schedulers
 - First come, first served
 - * Jobs in order they arrive
 - * Simple
 - * Problem: long jobs delay those behind
 - Shortest First
 - * Short jobs finish first, so long jobs do not cause delays
 - * Shortest remaining time first
 - reevaluate when new job is finished
 - * Problem
 - How to know how long a job is
 - Three level scheduling
 - * Jobs in input queue until moved to mem
 - Pick complementary Jobs: small/large; cpu heavy/io heavy
 - * CPU scheduler picks next job
 - * mem sched picks jobs to move from mem if out of memory
- Interactive Schedulers
 - Round Robin
 - * Each process gets fixed time slot (quantum)
 - * Rotate through ready processes
 - * Each makes progress
 - * What is a good amount of time?
 - Too short: Wasteful
 - Too long: not good for interactive
 - Typically: 10-100ms
 - Priority Scheduling
 - * Assign priority to each process
 - “Ready” process with highest priority runs next
 - * Dynamic priorities

- Reduce when using lots of CPU time
 - Increase on wait for I/O
- * Often grouped into multiple queues based on priority
 - Round robin within queues
- Shortest process next
 - * Run process that will finish soonest
 - * Guess time based on previous runs
 - * Not used because round robin works so well
- Lottery Scheduling
 - * Each process gets tickets for CPU time
 - More tickets -> More time
 - * Each quantum, a random ticket is picked
 - Can be implemented efficiently by using integer ranges for tickets
 - * Over time each process gets CPU m/n of the time
 - m = # of tickets process has
 - n = total # of tickets
 - * Tickets can be transferred
 - Processes exchange tickets
 - Clients to server
 - Parents to children

Scheduling in FreeBSD

Called ULE scheduler because file is sched_ule.c

- Two Parts
 - Low level
 - * Selects next thread from run queue
 - * Very fast
 - High level
 - * Sets thread priority
 - * Selects processor to use
 - * Slower
 - * Runs infrequently
 - * Adjusts quantum and priority
- Run queues only have ready threads
 - Blocked threads placed in
 - * Turnstile: short-term
 - * sleep queue: med/long term
 - Runnable thread that consumes quantum placed at end of queue

Policy VS Mechanism

- Separate what must be done (policy) from how it is done (mechanism)