

Week 2 (Jan 18 - 22)

Deadlocks

Resources

- Usually limited
- Examples
 - Printers
 - Semaphores
 - Memory
 - Tables
- Processes need access to resources
- Two types of resources
 - Pre-emptive: Can be taken away with no ill effects
 - Non-preemptable: Will cause failure if taken away
- Sequence of events to use resource
 1. Request
 2. Use
 3. Release
- Can't use if request is denied
 - Options:
 - * Block and Wait
 - * Continue (if possible) without resource
 - * Process fails with error code
 - Some of these options can prevent deadlock

Deadlocks

- When do deadlocks happen?
 - Processes are granted exclusive access to resources
 - Each deadlocked process needs a resource held by another deadlocked process
- What is deadlock?
 - A set of processes is deadlocked if each process in the set is waiting for an event that only another process can make happen
 - Usually the event is a release of a resource
 - In deadlock, none of the processes can:
 - * Run
 - * Release resources
 - * Be awakened
- Four conditions for deadlock (Coffman Conditions)
 1. Mutual Exclusion
 - Each resource assigned to at most one process
 2. Hold and Wait
 - Processes holding resources can request more resources

- 3. No preemptive
 - Previously granted resources can not be taken away
- 4. Circular Wait
 - Circular chain of processes each waiting for a resource held by another process in the chain
- These 4 conditions are necessary and sufficient for deadlock

How OS deals with Deadlock

- Ostrich method (Ignore it)
- Detect and Recover
- Dynamically avoid
- Prevent

Ostrich Method (ignore)

- Reasonable if deadlocks rare
 - Cost of prevention is high
- Unix/Windows take this approach
 - Resources plentiful
 - Deadlocks rare
 - Typically handled by rebooting
- Tradeoff: Convenience vs Correctness

Avoid Deadlock

- Find ways to detect and prevent deadlocks
- Stop them from happening in the first place

Deadlock detection using graphs

- Hold requests in table
- Graph contains cycle => deadlock
- Use DFS to find cycles

Recover from deadlock

- Through Preemption
 - Take resources from another process
 - Depends on nature of resource and process
- Rollback
 - Checkpoint a process periodically
 - Use saved state to restart process in deadlock
 - May present problem if process affects outside world
- Killing process
 - Crude and Simple
 - Kill one process in deadlock
 - Other processes can get its resources
 - Pick process that can most easily be rerun

Prevent completely

- Ensure one of necessary and sufficient conditions not true
 - Mutual Exclusion
 - Circular wait
 - Hold and Wait
 - No Pre-emption
- Not always possible

Attacking Mutual Exclusion

- Some devices (printers) can be spooled
 - Only printer daemon accesses actual printer
 - Spool contents as file
- Principle
 - Avoid assigning resources unless necessary
 - As few processes as possible claim resource

Attacking Hold and Wait

- Require processes to request resources before starting
 - Never wait for what it needs
 - Problems
 - * Cannot always know what it needs
 - * Takes up resources other processes could be using
 - Variation
 - * Give up all resources before making a new request
 - Granted all previous resources as well as all new resources

Attacking No Pre-emption

- Usually not viable
- May work for some resources
 - Forcibly take away memory pages, suspending process
 - Process may resume with no ill effects

Attacking Circular Wait

- Assign an order to resources
- Always acquire resources in numerical order
 - Need to acquire all at once
- Prevents circular wait
 - Processes holding resource n can't wait for m ($m < n$)
- No way to complete cycle

What does FreeBSD do?

- What resources are of issue?
 - Locks and Semaphores
 - Physical resources usually not a concern
- Goal: Prevent deadlocks
 - Low cost
 - Simple
 - Flexible

Two Phase Locking

- Phase One
 - Process tries to lock all data it needs at one time
 - If needed data is locked, start over
- Phase Two
 - Perform updates
 - Release lock
- Similar to request all at once
- Used in databases
- Avoids hold and wait

Starvation

- Some schedulers may cause jobs to be postponed indefinitely
 - e.g. shortest first schedule may cause a long job to never complete even when not blocked
- Starvation can lead to deadlock
 - Starved process may hold resources
 - If not used/released in reasonable time, shortage may cause deadlock

Livelock

- Processes run, but make no progress
 - Example:
 - * P0 has A ; P1 has B
 - * P0 wants B ; P1 wants A
 - * Both realize deadlock
 - * P0 drop A ; P1 drop B
 - * P0 gets B ; P1 gets A
 - * ...
- Real life example
 - Ethernet transmission collisions
 - * On collision wait and try again